

Autonomous Driving Project Report

Group Member:

Shiqi Lin, U1002894184

Fanxuan Guo, U1003110288

Introduction:

In this project, we stepped closer to application of computer vision for industry. The goal is to take sets of images and correctly detect road and cars. By using technique of deep learning over computer vision, we could increase the accuracy of object detections, thus lower the risk of autonomous driving. If this is done correctly, it could help vehicle detects obstacles and navigate on the road.

We contributed evenly for all subtasks of this project, and we worked closer to research and make decisions together. Here is the rough division:

- Shiqi Lin
Computed disparity, depth and best-fit plane. Wrote visualization of points cloud for 3D locations and plane. First method 1 of road detection classifier. Subtask 1, 2, 3, 4 and 5.
- Fanxuan Guo
Wrote method 2 of road detection classifier, car detection, car viewpoint classifier and visualization of car viewpoint prediction. Subtask 3, 6, 7 and 8.

Subtask 1: Compute Disparity (*image_processing.py*)

For this task, we implement 2 methods.

- Method 1: *compute_disparity()* based on the logic from lecture 12

For each pixel in the left image, we create a patch and compare it with the patches on the right stereo image on the same row. The compared pixel locations are from $\max(0, \text{left_pixel_x} - 16)$ to left_pixel_x (same row, diff colm index). We choose 16 as max disparity to increase speed. We also create a matrix to store the vectorized patches that were seen in the right image for fast retrieval.



Figure 1: Method1 Disparity of *test/image_left/umm_000087.jpg*

- Method 2: *compute_disparity_by_cv2()*

We use opencv library to compute disparity. In order to create a smoother result, we smooth the disparity output by a 5 x 5 gaussian filter.




Figure 2: Method2 Disparity of *test/image_left/umm_000087.jpg*

- Comparison:

Our own method has a reasonable running time when patch size is small, but when it's too large, the method is too slow. However, method 2 runs fast under any condition.

Subtask 2: Compute Depth (*image_processings.py*)

This Task we use the code from A4 to generate the depth map.

Formula from the lecture: $Depth = \frac{\text{focal length} * \text{baseline}}{\text{disparity}}$




Figure 3: Method2 Disparity -> Depth Map of *test/image_left/umm_000087.jpg*



Figure 4: Method2 Disparity -> Depth Map of *test/image_left/umm_000087.jpg*

Subtask 3: Road Detection (*road_detection.py*, *road_segmentation.py*)

- Method 1:

We use the CNN-LSTM network structure proposed from Road Segmentation Using CNN and Distributed LSTM (<https://arxiv.org/pdf/1808.04450.pdf>)

Based on the pyramid prediction scheme, we firstly crop (at the center) and scale the input image to size of 600×160 . In the image pre-processing stage, for each image, we add the x, y pixel location to the RGB channel and form a 5 channels image tensor. This way we can not only extract the color, edge, point features using NN but also predict result based on the pixel location as in the KITTI dataset, the road always appears on the bottom of the picture. We also processed each gt_mask to generate a ground truth boundary vector used in the loss calculation stage. And then following the proposed network structure, we build our own torch network model with learning rate $1e-5$, Adam optimizer, L1 loss function and 80 epochs. The proposed network structure is as following: (picture from the paper) The network contains an encoder, feature processor and a decoder. It finally outputs a boundary vector of the road.

Layer	Kernel ($w_1 \times h_1 \times d_1$)	Input ($w \times h \times d$)	Output ($w_2 \times h_2 \times d_1$)
Input	—	$600 \times 160 \times 5$	—
Conv1 s=2	$3 \times 3 \times 64$	$600 \times 160 \times 5$	$600 \times 160 \times 64$
Conv2	$3 \times 3 \times 64$	$600 \times 160 \times 64$	$600 \times 160 \times 64$
Conv3 s=2	$3 \times 3 \times 64$	$300 \times 80 \times 64$	$300 \times 80 \times 64$
Conv4	$3 \times 3 \times 64$	$300 \times 80 \times 64$	$300 \times 80 \times 64$
Conv5 s=2	$3 \times 3 \times 64$	$150 \times 40 \times 64$	$150 \times 40 \times 64$
Conv6	$3 \times 3 \times 64$	$150 \times 40 \times 64$	$150 \times 40 \times 64$
Conv7	$3 \times 3 \times 64$	$75 \times 20 \times 64$	$75 \times 20 \times 64$
D-LSTM 1	64	$75 \times 20 \times 64$	$75 \times 20 \times 64$
Conv8	$3 \times 3 \times 64$	$75 \times 20 \times 64$	$75 \times 20 \times 64$
D-LSTM 2	64	$75 \times 20 \times 64$	$75 \times 20 \times 64$
Conv9	$1 \times 5 \times 64$	$75 \times 20 \times 64$	$75 \times 4 \times 64$
Conv10	$1 \times 4 \times 1$	$75 \times 4 \times 64$	$75 \times 1 \times 1$
Up- sample	$8 \times 1 \times 1$	$75 \times 1 \times 1$	$600 \times 1 \times 1$
Output	—	—	$600 \times 1 \times 1$

Figure 5: Detailed Network Blocks

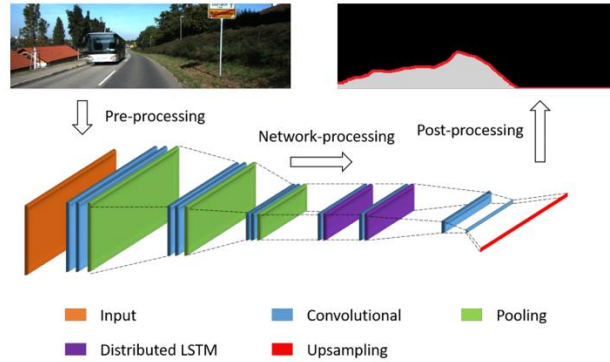


Figure 6: Proposed Network Diagram

As the paper proposed, this network should produce a result with a high accuracy rate ($\sim 89\%$) within a reasonably short time. However, during real implementation, we found out that the model actually takes a long time to train.

As we take the pixel location as a road feature, the training data bias influences the final prediction a lot. Since in many images, the road appears on the bottom left corner, our network predicts that the road would appear on the left bottom more possibly. This leads to the false positive of the road prediction when the road appears on the right side of the image.




Figure 7 (top) and 8 (bottom): Bad Example of CNN-LSTM Road Classifier for
test/image_left/umm_000087.jpg

- Method 2:

We chose the U-Net convolutional network. The U-Net is designed for image segmentation. The structure contains two parts -- encoder and decoder. The encoder is used to capture image features, and the decoder is used to perform localization. For our task, we need to pass in an image as training data and classify a ground truth mask for the road. Since U-Net is an end-to-end fully convolutional network, and we need the input shape and output shape to be the same, so U-Net is a good choice.

We found the detailed structure and implementation of U-Net NN in here:

<https://www.depends-on-the-definition.com/unet-keras-segmenting-images/>.

For the parameters of U-Net, we chose batch size to be 20, and split 80% of data to be training data set, rest of them to be validation data set. Figure 9 and Figure 10 are the loss curve and Accuracy curve during the training process. We could see that the validation accuracy is about 80%.




Figure 9 (left) and 10 (right): Accuracy and Loss plot for model

In Figure 13, 14, 17 and 18, we plot the ground truth mask returned from the U-Net and plot the segmentation on the image. Since the KITTI ground truth mask only covers the driving lane that the car is currently on, so most of the prediction mask gets the middle road.

- Comparison:

As the first method CNN-LSTM we tried took a very long time to train and the result is biased based on the road location. We tried the second method U-Net which is simpler and the training is faster. After comparison, we found out that although method 2 creates a smoother boundary for the road segmentation, the prediction result is unstable.




Figure 11 (top) and 12 (bottom): CNN-LSTM Road Classifier for *test/image_left/umm_000073.jpg*




Figure 13 (top) and 14 (bottom): U-Net Road Classifier for *test/image_left/umm_000073.jpg*




Figure 15 (top) and 16 (bottom): CNN-LSTM Road Classifier for *test/image_left/umm_000087.jpg*




Figure 17 (top) and 18 (bottom): U-Net Road Classifier for *test/image_left/umm_000087.jpg*

Subtask 4: Fit Ground Plane (*road_visualization.py*)

This task requires depth map and ground truth mask from previous Subtask. We first computed 3D location based on depth map. We used the formula $X = \frac{Z-(x*Px)}{f}$, $Y = \frac{Z-(y*Py)}{f}$, where Z is the depth in the depth map, x and y is the image coordinate. And we found Px, Py and focal length from calibration information.

Then for all the 3D location, we used ground truth mask to filter out points that are not road. The last step is to fit a plane. Here we used a technique called Least Square to find the best fit plane. The algorithm can find a plane that minimize the least square distance to every point. In order for our algorithm to be robust to outliers, we applied RANSAC over Least Square. We had 3000 iterations, each iteration we took 5 points and got a best fit plane using Least Square. Then we computed outliers with threshold 0.02. After all the iterations, we used the plane that had the least number of outliers.

Figure 19 used matplotlib to plot all the road 3D points and the best fit plane after RANSAC. We can tell that the plane is almost touching the bottom side of points cloud from Figure 20.

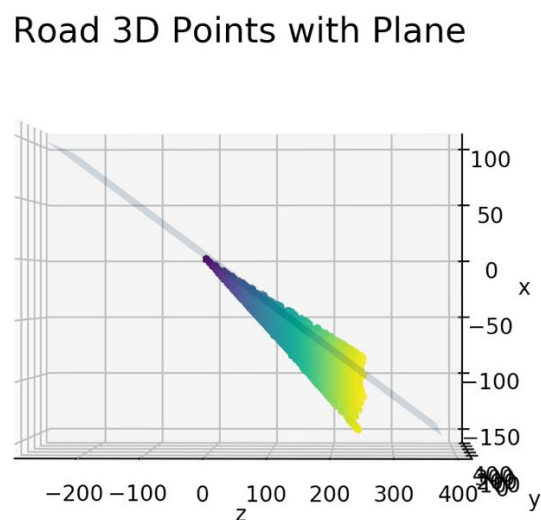
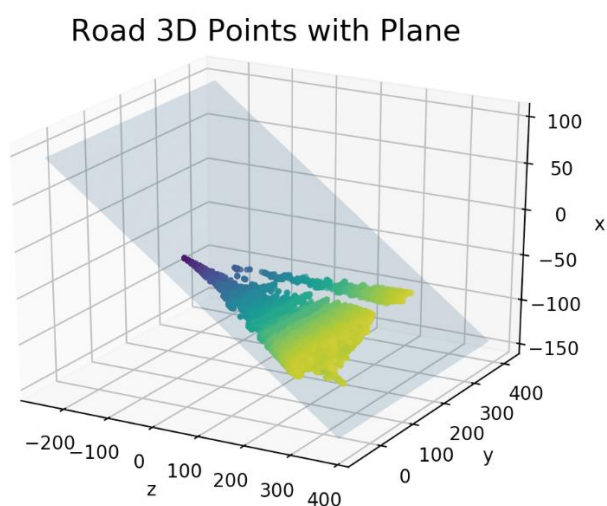


Figure 19 (left) and 20 (right): Road Plane for *test/image_left/umm_000087.jpg*

Subtask 5: 3D Point Cloud (*road_visualization.py*)

This Task requires best fit plane and 3D locations from previous Subtask. We used open3D to put the color onto 3D location points cloud, then we generated another 3D points cloud for the plane and add it to the figure.

Here is the image from *test/image_left/umm_000087.jpg* that we used to plot 3D points cloud and ground plane.



Figure 21 is captured when we rotate the model on the side. Figure 22 is the front face, and Figure 23 is viewed from top right.




Figure 21 (left) and 22 (right): 3D Point Cloud for *test/image_left/umm_000087.jpg*




Figure 23: 3D Point Cloud with OpenCV for *test/image_left/umm_000087.jpg*

Subtask 6: Car Detection (*car_detection.py*)

At the beginning we searched a lot of pre-trained model for KITTI data set, but they have low accuracy on the test image. Finally, we used Faster RCNN from torchvision pre-trained model. It is not designed for car detection, we filtered out other objects (E.g. person). The only objects we used to detect are car, bus and truck.

The reason we chose Faster RCNN is because of its high accuracy and short computation time, compared with YOLO and SSD. YOLO and SSD have better performance for object detection in video. Here is the reference code of how to use torchvision pre-trained model:

<https://www.learnopencv.com/faster-r-cnn-object-detection-with-pytorch/>



Figure 24: Car Detection for *train/image_left/um_000010.jpg*

Subtask 7: Viewpoint Classifier (*car_viewpoint.py*)

For this task, we built a simple CNN using Keras to do viewpoint classification. We used the one-hot encoder to represents the training labels. There are 12 classes, so each label has dimension 12 x 1. Figure 25 is the structure of our CNN.

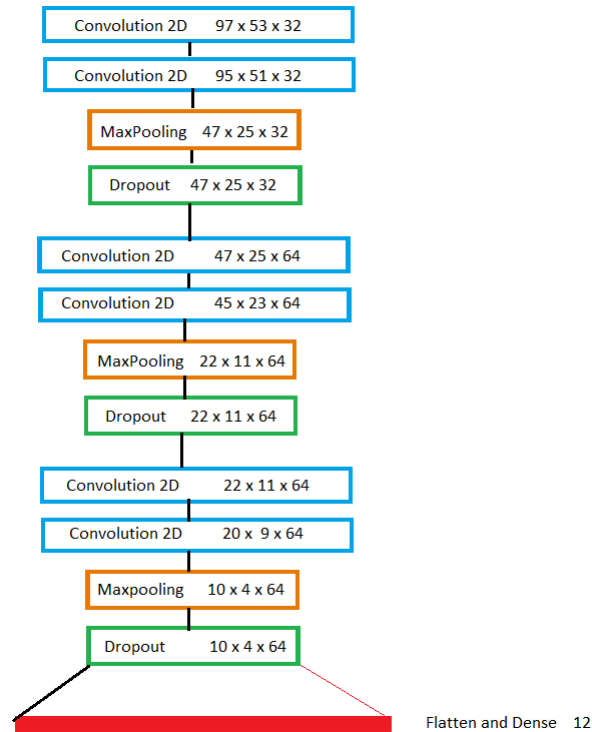



Figure 25: Structure of CNN model

Here is the reference code of simple CNN using Keras: <https://www.learnopencv.com/image-classification-using-convolutional-neural-networks-in-keras/>. We could see that the validation accuracy is 94%, but the training accuracy is higher than validation accuracy. The potential reason for this is that our model is overfitting.



Subtask 8: Car and Viewpoint (*car_visualization.py*)

This task requires pre-trained viewpoint model and car detection. For each image, we first detected cars and ignore bounding box that are too small. Then we pass-in the car image patch to the pre-trained model and plot the viewpoint angle with arrows. Arrow pointing down is 0° , and arrow pointing left is 90° .




Figure 28: Car and Viewpoint for *train/image_left/um_000010.jpg*

image_processings.py

```
1 import numpy as np
2 import cv2
3 from matplotlib import pyplot as plt
4 import skimage.io
5
6
7 def compute_disparity(left, right, p):
8     """
9     Project 2: Question 1
10    Method 1
11
12    :param left: the left image in the stereo image pairs
13    :type left: numpy array
14    :param right: the right image in the stereo image pair
15    :type right: numpy array
16    :param p: patch size
17    :type p: int
18    :return: disparity map of the image pair
19    :rtype: numpy array of the same shape as the input images
20    """
21    # max disparity to check
22    max_d = 16
23
24    disp_map = np.zeros(left.shape)
25
26    # sizes
27    im_height, im_width = left.shape
28    p_half = p // 2
29
30    # zero-padding: N + p - 1 * M + p - 1 for left and right image
31    padded_shape = im_height + p - 1, im_width + p - 1
32
33    padded_im_left = np.zeros(padded_shape)
34    padded_im_left[p_half:(im_height + p_half), p_half:(im_width + p_half)] = left
35
36    padded_im_right = np.zeros(padded_shape)
37    padded_im_right[p_half:(im_height + p_half), p_half:(im_width + p_half)] = right
```



```

38
39 # matrix of vectorized patches for right image for fast retrieval
40 r_patch_vectorized = np.zeros((im_height, im_width, p * p))
41
42 for i in range(im_height):
43     for j in range(im_width):
44         l_patch = padded_im_left[i: i + p, j: j + p].reshape(-1)
45         r_patch = padded_im_right[i: i + p, j: j + p].reshape(-1)
46
47         r_patch_vectorized[i, j, :] = r_patch
48
49         # calculate absolute intensity difference
50         start = max(0, j - max_d)
51         abs_diff = np.abs(np.copy(r_patch_vectorized[i, start: j + 1, :]) - l_patch)
52         ssd = np.sum(abs_diff, axis=1)
53
54         # find the best match
55         ssd = np.flip(ssd, 0)
56         min_d = np.argmin(ssd)
57
58         disp_map[i, j] = min_d
59
60 cv2.normalize(disp_map, disp_map, 0, 255, cv2.NORM_MINMAX)
61 return disp_map

```

```

64 def compute_disparity_by_cv2(left, right, block_size):
65     """
66     Project 2: Question 1
67     Method 2
68
69     :param left: the left image in the stereo image pairs
70     :type left: numpy array
71     :param right: the right image in the stereo image pair
72     :type right: numpy array
73     :return: disparity map of the image pair
74     :rtype: numpy array of the same shape as the input images
75     """
76     stereoMatcher = cv2.StereoBM_create()
77     stereoMatcher.setMinDisparity(4)
78     stereoMatcher.setNumDisparities(128)
79     stereoMatcher.setBlockSize(block_size)
80     stereoMatcher.setSpeckleRange(16)
81     stereoMatcher.setSpeckleWindowSize(10)
82     disp = stereoMatcher.compute(left, right)
83
84     kernel = np.ones((5, 5), np.float32) / 25
85     disp = cv2.filter2D(disp, -1, kernel)
86     return disp

```

```

89 def compute_depth(disp_map, f, t):
90     """
91     Project 2: Question 2
92
93     :param disp_map: disparity map of the stereo image pair
94     :type disp_map: numpy array
95     :param f: focal length
96     :type f: float
97     :param t: baseline
98     :type t: float
99     :return: depth map of the image pair
100     :rtype: numpy array of the same shape as the input disp_map
101     """
102     depth = np.divide(-f * t, disp_map, where=disp_map != 0)
103     cv2.normalize(depth, depth, 0, 255, cv2.NORM_MINMAX)
104     return depth

```

```

107 if __name__ == "__main__":
108     F = 7.215377000000e+02
109     T = 0.54 * 1000
110     PX, PY = 6.095593000000e+02, 1.728540000000e+02
111
112     left_img = cv2.imread('./data/test/image_left/umm_000087.jpg', cv2.IMREAD_GRAYSCALE)
113     right_img = cv2.imread('./data/test/image_right/umm_000087.jpg', cv2.IMREAD_GRAYSCALE)
114
115     # Q1: Compute disparity
116     # Method 1
117     # disp = compute_disparity(left_img, right_img, 10)
118     # Method 2
119     disp = compute_disparity_by_cv2(left_img, right_img, 5)
120
121     plt.imshow(disp, 'gray')
122     plt.title('Disparity', fontsize=16)
123     plt.show()
124
125     # Q2: Compute depth
126     depth = compute_depth(disp, F, T)
127     plt.imshow(depth, 'gray')
128     plt.title('Depth', fontsize=16)
129     plt.show()
130     skimage.io.imsave('./Output Images/depth_umm_000087.png', depth.astype('uint8'))
131

```

road_segmentation.py (Method 1)

```

1  # ref: https://arxiv.org/pdf/1808.04450.pdf
2  # Method 1 for Q3
3  import ...
17
18  # dataloader
19  class KittiDataset(tu.data.Dataset):
20      def __init__(self, image_path, label_path, transform):
21          self.imgs, self.labels = read_files(image_path, label_path)
22          self.transform = transform
23
24      def __getitem__(self, index):
25          img = self.imgs[index]
26          label = self.labels[index]
27          img = self.transform(img)
28
29          dt = {'img': img,
30               'gt': torch.IntTensor(label)}
31
32          return dt
33
34      def __len__(self):
35          return len(self.imgs)
36
37      def read_files(image_path, label_path):
38          # read image and pre-process them for the training data loader
39
40          images = []
41          labels = []
42
43          for img_name in os.listdir(image_path):
44              path = os.path.join(image_path, img_name)
45              if "DS_Store" not in path:
46
47                  # process image with pyramid prediction scheme
48                  img = Image.open(path).convert('RGB')
49                  # crop image at center
50                  c_img = crop_image(img)

```

```

51 # rescale image to the same size
52 img = ImageOps.fit(img, (600, 160), Image.ANTIALIAS)
53 # add x, y coordinate channels
54 img = generate_coord_channels(img)
55 c_img = generate_coord_channels(c_img)
56 images.append(img)
57 images.append(c_img)
58
59 # process gt_mask with pyramid prediction scheme
60 # get gt_mask name
61 img_name = img_name.split('.')[0]
62 img_name_l = img_name.split('_')
63 img_name_l.insert(1, "road")
64 gt_name = '_'.join(img_name_l)
65 gt_name += '.png'
66
67 gt_path = os.path.join(label_path, gt_name)
68 gt_mask = Image.open(path).convert('L')
69
70 c_mask = crop_image(gt_mask)
71 gt_mask = ImageOps.fit(gt_mask, (600, 160), Image.ANTIALIAS)
72
73 # generate ground truth road segmentation boundary vector
74 # vector size = 600 * 1, v[i] = the row index of the boundary pixel
75 gt_mask = np.array(gt_mask)
76 bound = gt_mask.argmax(axis=0)
77 gt_col_sum = np.sum(gt_mask, axis=0)
78 bound[gt_col_sum == 0] = 160
79
80 c_mask = np.array(c_mask)
81 c_bound = c_mask.argmax(axis=0)
82 c_col_sum = np.sum(c_mask, axis=0)
83 c_bound[c_col_sum == 0] = 160
84
85 labels.append(bound)

```

```

86         labels.append(c_bound)
87
88     return images, labels
89
90 def make_coordinates_matrix(im_shape, step=1):
91     """
92     Return a matrix of size (im_shape[0] x im_shape[1] x 2) such that g(y,x)=[y,x]
93     """
94     range_x = np.arange(0, im_shape[1], step)
95     range_y = np.arange(0, im_shape[0], step)
96     axis_x = np.repeat(range_x[np.newaxis, ...], len(range_y), axis=0)
97     axis_y = np.repeat(range_y[... , np.newaxis], len(range_x), axis=1)
98
99     return np.dstack((axis_y, axis_x))
100
101 def generate_coord_channels(image):
102     np_img = np.array(image)
103     coord_m = make_coordinates_matrix(np_img.shape)
104     np_img = np.concatenate((np_img, coord_m), axis=2)
105     return np_img
106
107 def crop_image(image):
108     center_x = image.size[0] / 2
109     center_y = image.size[1] / 2
110     left = center_x - 300
111     top = center_y - 80
112     right = center_x + 300
113     bottom = center_y + 80
114     c_img = image.crop((left, top, right, bottom))
115     return c_img

```

```

117 # training network
118 class CNN_LSTM(nn.Module):
119     def __init__(self):
120         super(CNN_LSTM, self).__init__()
121         self.conv_in = nn.Conv2d(in_channels=5, out_channels=64, stride=1, kernel_size=3, padding=1) # 2
122         self.conv_en = nn.Conv2d(in_channels=64, out_channels=64, stride=1, kernel_size=3, padding=1) # 2
123         self.pool = nn.MaxPool2d(2, 2)
124
125         self.conv_fp = nn.Conv2d(in_channels=64, out_channels=64, stride=1, kernel_size=3, padding=1)
126         self.lstm = nn.LSTM(input_size=64, hidden_size=64)
127
128         self.conv_de1 = nn.Conv2d(in_channels=64, out_channels=64, stride=(5, 1), kernel_size=(5, 1))
129         self.conv_de2 = nn.Conv2d(in_channels=64, out_channels=1, stride=1, kernel_size=(4, 1))
130         self.upsample = nn.Upsample(scale_factor=(1, 8), mode='nearest')
131
132     def forward(self, image):
133         # print("input shape: expect shape 5 * 160 * 600 ",image.shape)
134         feature = self.conv_in(image)
135         # print("conv1: expect shape 64 * 160 * 600 ",feature.shape)
136         feature = self.conv_en(feature)
137         # print("conv2: expect shape 64 * 160 * 600 ",feature.shape)
138         feature = self.pool(feature)
139         # print("pool1: expect shape 64 * 80 * 300 ",feature.shape)
140
141         feature = self.conv_en(feature)
142         # print("conv3: expect shape 64 * 80 * 300 ",feature.shape)
143         feature = self.conv_en(feature)
144         # print("conv4: expect shape 64 * 80 * 300 ",feature.shape)
145         feature = self.pool(feature)
146         # print("pool2: expect shape 64 * 40 * 150 ",feature.shape)
147
148         feature = self.conv_en(feature)
149         # print("conv5: expect shape 64 * 40 * 150 ",feature.shape)
150         feature = self.conv_en(feature)
151         # print("conv6: expect shape 64 * 40 * 150 ",feature.shape)
152         feature = self.pool(feature)
153
154         output = self.conv_fp(feature)
155         # print("conv7: expect shape 64 * 20 * 75 ", output.shape)
156         output = output.view(output.shape[2], output.shape[3], 64)
157         hidden_state = torch.randn(1, output.shape[1], 64)
158         cell_state = torch.randn(1, output.shape[1], 64)
159         output = self.lstm(output, (hidden_state, cell_state))[0]
160         output = output.view(1, 64, output.shape[0], output.shape[1])
161         # print("lstm1: expect shape 64 * 20 * 75 ", output.shape)
162
163         output = self.conv_fp(output)
164         # print("conv8: expect shape 64 * 20 * 75 ", output.shape)
165         output = output.view(output.shape[2], output.shape[3], 64)
166         hidden_state = torch.randn(1, output.shape[1], 64)
167         cell_state = torch.randn(1, output.shape[1], 64)
168         output = self.lstm(output, (hidden_state, cell_state))[0]
169         output = output.view(1, 64, output.shape[0], output.shape[1])
170         # print("lstm2: expect shape 64 * 20 * 75 ", output.shape)
171
172         output = self.conv_de1(output)
173         # print("conv9: expect shape 64 * 4 * 75 ", output.shape)
174         output = self.conv_de2(output)
175         # print("conv10: expect shape 1 * 1 * 75 ", output.shape)
176         output = self.upsample(output)
177         # print("final output: expect shape 1 * 1 * 600 ", output.shape)
178         return output

```

```

181 def train_road_classifier(md, optimizer, trainloader, criterion):
182     num_epochs = 80
183     total_step = 270
184     losses = list()
185
186     for epoch in range(1, num_epochs + 1):
187         running_loss = 0.0
188         for i, data in enumerate(trainloader, 0):
189             # zero the gradients
190             md.zero_grad()
191             optimizer.zero_grad()
192
193             # set model into train mode
194             md.train()
195
196             images = data['img']
197             gt_boundaries = data['gt']
198
199             # Pass the inputs through the CNN-RNN model.
200             outputs = md(images.float())
201
202             # Calculate the batch loss
203             loss = criterion(outputs[0, 0, :, :], gt_boundaries.float())
204
205             # Backward pass
206             loss.backward()
207
208             # Update the parameters in the optimizer
209             optimizer.step()
210
211             losses.append(loss.item())
212             running_loss += loss.item()
213
214             # save the losses
215             np.save('losses', np.array(losses))

```



```

217         # Get training statistics.
218         stats = 'Epoch [%d/%d], Step [%d/%d], Loss: %.4f' % (
219             epoch, num_epochs, i, total_step, loss.item())
220
221         # Print training statistics (on same line).
222         print('\n' + stats, end="")
223         sys.stdout.flush()
224
225     # Save the weights.
226     print("\nSaving the model")
227     torch.save(md.state_dict(), os.path.join('./data/train/results-dets', 'md.pth'))
228
229     # testing visualization
230     def visualize_segmentation(boundary):
231         # visualize the segmentation mask for the image given the boundary vector
232         image = Image.new("RGB", (600, 160))
233         draw = ImageDraw.Draw(image)
234         points = generate_boundary_coord(boundary)
235         points = np.vstack((points, points[0, :]))
236         points = tuple(map(tuple, points))
237         draw.polygon((points), fill=(255, 255, 255))
238         # image.show()
239         return image
240
241     def generate_boundary_coord(y, step=1):
242         # generate the road boundary xy coordinate based on the boundary row index vector
243         im_shape = (1, y.shape[1])
244
245         range_x = np.arange(0, im_shape[1], step)
246         range_y = np.arange(0, im_shape[0], step)
247         axis_x = np.repeat(range_x[np.newaxis, ...], len(range_y), axis=0)
248         result = np.dstack((axis_x, y)).reshape((im_shape[1], 2))
249
250         return result
251
252     def reassemble(scaled_img, cropped_img, img_size):
253         # reassemble the overlayed scaled and cropped images
254         center_x = img_size[1] // 2
255         center_y = img_size[0] // 2
256         # scale the image back to its original size
257         scaled_img = ImageOps.fit(scaled_img, img_size, Image.ANTIALIAS)
258         # replaced the pixel values with the cropped image at the center
259         scaled_img.paste(cropped_img, (center_x, center_y))
260         # scaled_img.show()
261         return scaled_img

```

```

263 def overlay_mask(image, gt_mask):
264     gt_mask = gt_mask.astype(int)
265     # code from tutorial7 to visualize segmentation
266     obj_ids = np.unique(gt_mask)
267     number_object = obj_ids.shape[0]
268
269
270     count = 0
271     for o_id in obj_ids:
272         gt_mask[gt_mask == o_id] = count
273         count += 1
274
275     base_COLORS = []
276
277     for key, value in mcolors.CSS4_COLORS.items():
278         rgb = webcolors.hex_to_rgb(value)
279         base_COLORS.append([rgb.blue, rgb.green, rgb.red])
280     base_COLORS = np.array(base_COLORS)
281
282     np.random.seed(99)
283     base_COLORS = np.random.permutation(base_COLORS)
284
285     colour_id = np.array([(id % len(base_COLORS) for id in range(number_object))])
286     COLORS = base_COLORS[colour_id]
287     COLORS = np.vstack([[0, 0, 0], COLORS]).astype("uint8")
288     mask = COLORS[gt_mask]
289     output = ((0.4 * image) + (0.6 * mask)).astype("uint8")
290     plt.imshow(output[:, :, ::-1])
291     plt.show()

```

```

293 if __name__ == "__main__":
294     # generate dataset loader
295     transform = transforms.Compose([transforms.ToTensor()])
296
297     # ===== comment out this part if you just want to see the test result =====
298     kd = KittiDataset(image_path='./data/train/image_left', label_path='./data/train/gt_image_left', transform=transform)
299
300     # training
301     model = CNN_LSTM().float()
302     optimizer = optim.Adam(model.parameters(), lr=1e-5)
303     criterion = nn.L1Loss()
304     trainloader = tu.data.DataLoader(kd, batch_size=1, shuffle=False, num_workers=2)
305     train_road_classifier(model, optimizer, trainloader, criterion)
306     # =====
307
308     # testing
309     md = CNN_LSTM().float()
310     cmd = CNN_LSTM().float()
311
312     md.load_state_dict(torch.load(os.path.join('./data/train/results-dets', 'md.pth')))
313     cmd.load_state_dict(torch.load(os.path.join('./data/train/results-dets', 'md.pth')))
314
315     image = Image.open('./data/test/image_left/uu_000041.jpg').convert('RGB')
316     img_s = image.size
317
318     # crop image at center with window size 600 * 160
319     c_img = crop_image(image)
320     c_img = generate_coord_channels(c_img)
321
322     # scale original image to 600 * 160
323     image = ImageOps.fit(image, (600, 160), Image.ANTIALIAS)
324     image = generate_coord_channels(image)

```

```

326     # turn image into torch tensor for prediction
327     image = transform(image).view(1, 5, 160, 600)
328     c_img = transform(c_img).view(1, 5, 160, 600)
329
330     # predict
331     outputs = md(image.float())
332     outputs = outputs.detach().numpy().reshape((1, 600))
333
334     coutputs = cmd(c_img.float())
335     coutputs = coutputs.detach().numpy().reshape((1, 600))
336
337     sv = visualize_segmentation(outputs)
338     cv = visualize_segmentation(coutputs)
339
340     final = np.array(reassemble(sv, cv, img_s))[:, :, 0]
341     overlay_mask(cv2.imread('./data/test/image_left/umm_000087.jpg'), final)

```

road_detection.py (Method 2)

```
1  # ref: https://www.depends-on-the-definition.com/unet-keras-segmenting-images/
2
3  import os
4  import cv2
5  import numpy as np
6  from keras import Model
7  from keras.callbacks import EarlyStopping, ReduceLROnPlateau, ModelCheckpoint
8  from keras.engine.saving import load_model
9  from keras.layers import Conv2D, BatchNormalization, Activation, MaxPooling2D, Dropout, Conv2DTranspose, concatenate
10 from skimage import io
11 import matplotlib.pyplot as plt
12 import matplotlib.colors as mcolors
13 import webcolors
14
15 im_width = 512
16 im_height = 128
17 split = 0.8
18
19 callbacks = [
20     EarlyStopping(patience=10, verbose=1),
21     ReduceLROnPlateau(factor=0.1, patience=3, min_lr=0.00001, verbose=1),
22     ModelCheckpoint('road_model_checkpoint.h5', verbose=1, save_best_only=True, save_weights_only=True)
23 ]
24
25
26 def conv2d_block(input_tensor, n_filters, kernel_size=3, batchnorm=True):
27     # first layer
28     x = Conv2D(filters=n_filters, kernel_size=(kernel_size, kernel_size), kernel_initializer="he_normal",
29               padding="same")(input_tensor)
30     if batchnorm:
31         x = BatchNormalization()(x)
32     x = Activation("relu")(x)
33
34     # second layer
35     x = Conv2D(filters=n_filters, kernel_size=(kernel_size, kernel_size), kernel_initializer="he_normal",
36               padding="same")(x)
37     if batchnorm:
38         x = BatchNormalization()(x)
39     x = Activation("relu")(x)
40     return x
```

```

43 def get_unet(input_img, n_filters=16, dropout=0.5, batchnorm=True):
44     # contracting path -- encoder
45     c1 = conv2d_block(input_img, n_filters=n_filters * 1, kernel_size=3, batchnorm=batchnorm)
46     p1 = MaxPooling2D((2, 2))(c1)
47     p1 = Dropout(dropout * 0.5)(p1)
48
49     c2 = conv2d_block(p1, n_filters=n_filters * 2, kernel_size=3, batchnorm=batchnorm)
50     p2 = MaxPooling2D((2, 2))(c2)
51     p2 = Dropout(dropout)(p2)
52
53     c3 = conv2d_block(p2, n_filters=n_filters * 4, kernel_size=3, batchnorm=batchnorm)
54     p3 = MaxPooling2D((2, 2))(c3)
55     p3 = Dropout(dropout)(p3)
56
57     c4 = conv2d_block(p3, n_filters=n_filters * 8, kernel_size=3, batchnorm=batchnorm)
58     p4 = MaxPooling2D(pool_size=(2, 2))(c4)
59     p4 = Dropout(dropout)(p4)
60
61     c5 = conv2d_block(p4, n_filters=n_filters * 16, kernel_size=3, batchnorm=batchnorm)
62
63     # expansive path -- decoder
64     u6 = Conv2DTranspose(n_filters * 8, (3, 3), strides=(2, 2), padding='same')(c5)
65     u6 = concatenate([u6, c4])
66     u6 = Dropout(dropout)(u6)
67     c6 = conv2d_block(u6, n_filters=n_filters * 8, kernel_size=3, batchnorm=batchnorm)
68
69     u7 = Conv2DTranspose(n_filters * 4, (3, 3), strides=(2, 2), padding='same')(c6)
70     u7 = concatenate([u7, c3])
71     u7 = Dropout(dropout)(u7)
72     c7 = conv2d_block(u7, n_filters=n_filters * 4, kernel_size=3, batchnorm=batchnorm)
73
74     u8 = Conv2DTranspose(n_filters * 2, (3, 3), strides=(2, 2), padding='same')(c7)
75     u8 = concatenate([u8, c2])
76     u8 = Dropout(dropout)(u8)
77     c8 = conv2d_block(u8, n_filters=n_filters * 2, kernel_size=3, batchnorm=batchnorm)
78
79     u9 = Conv2DTranspose(n_filters * 1, (3, 3), strides=(2, 2), padding='same')(c8)
80     u9 = concatenate([u9, c1], axis=3)
81     u9 = Dropout(dropout)(u9)
82     c9 = conv2d_block(u9, n_filters=n_filters * 1, kernel_size=3, batchnorm=batchnorm)
83
84     outputs = Conv2D(1, (1, 1), activation='sigmoid')(c9)
85     model = Model(inputs=[input_img], outputs=[outputs])
86     return model

```

```

89 def overlay_mask(image, mask):
90     # pre-process mask
91     mask[np.where(mask < 0.5)] = 0 # not road pixel
92     mask[np.where(mask >= 0.5)] = 1 # road pixel
93     mask = mask.astype(int)
94
95     # code from tutorial7 to visualize segmentation
96     obj_ids = np.unique(mask)
97     number_object = obj_ids.shape[0]
98
99     count = 0
100    for o_id in obj_ids:
101        mask[mask == o_id] = count
102        count += 1
103
104    base_COLORS = []
105
106    for key, value in mcolors.CSS4_COLORS.items():
107        rgb = webcolors.hex_to_rgb(value)
108        base_COLORS.append([rgb.blue, rgb.green, rgb.red])
109    base_COLORS = np.array(base_COLORS)
110
111    np.random.seed(99)
112    base_COLORS = np.random.permutation(base_COLORS)
113
114    colour_id = np.array([(id % len(base_COLORS) for id in range(number_object))])
115    COLORS = base_COLORS[colour_id]
116    COLORS = np.vstack([[0, 0, 0], COLORS]).astype("uint8")
117    mask = COLORS[mask]
118    output = ((0.4 * image) + (0.6 * mask)).astype("uint8")
119    io.imshow(output)
120    io.show()

```



```

123 def visualize_performance(history):
124     # Loss Curves
125     plt.figure(figsize=[8, 6])
126     plt.plot(history.history['loss'], 'r', linewidth=3.0)
127     plt.plot(history.history['val_loss'], 'b', linewidth=3.0)
128     plt.legend(['Training loss', 'Validation Loss'], fontsize=18)
129     plt.xlabel('Epochs ', fontsize=16)
130     plt.ylabel('Loss', fontsize=16)
131     plt.title('Loss Curves', fontsize=16)
132     plt.show()
133
134     # Accuracy Curves
135     plt.figure(figsize=[8, 6])
136     plt.plot(history.history['accuracy'], 'r', linewidth=3.0)
137     plt.plot(history.history['val_accuracy'], 'b', linewidth=3.0)
138     plt.legend(['Training Accuracy', 'Validation Accuracy'], fontsize=18)
139     plt.xlabel('Epochs ', fontsize=16)
140     plt.ylabel('Accuracy', fontsize=16)
141     plt.title('Accuracy Curves', fontsize=16)
142     plt.show()

```

```

145 def visualize_predict(model_name):
146     model = load_model(model_name)
147
148     images = []
149     images_path = []
150     images_shape = []
151     for path in os.listdir('data/test/image_left'):
152         img_full_path = 'data/test/image_left' + '/' + path
153         if "DS_Store" not in img_full_path:
154             img = io.imread(img_full_path, as_gray=True)
155             images_shape.append(img.shape)
156             img = cv2.resize(img, dsize=(img_width, img_height), interpolation=cv2.INTER_CUBIC)
157             img = img.reshape((img_height, img_width, 1))
158             images.append(img)
159             images_path.append(img_full_path)
160
161     prediction = model.predict(np.array(images), verbose=1)
162
163     for i in range(len(images)):
164         pred = prediction[i][:, :, 0]
165         pred = cv2.resize(pred, dsize=(images_shape[i][1], images_shape[i][0]), interpolation=cv2.INTER_CUBIC)
166         io.imshow(pred)
167         io.show()
168         img = io.imread(images_path[i])
169         overlay_mask(img, pred)

```

```

173 def load_data():
174     # Read images
175     images = []
176     masks = []
177     for path in os.listdir('data/train/image_left'):
178         img_full_path = 'data/train/image_left' + '/' + path
179         if "DS_Store" not in img_full_path:
180             img = io.imread(img_full_path, as_gray=True)
181             img = cv2.resize(img, dsize=(im_width, im_height), interpolation=cv2.INTER_CUBIC)
182             img = img.reshape((im_height, im_width, 1))
183             images.append(img)
184
185             num = path.split('.')[0].split('_')[1]
186             mask_full_path = 'data/train/gt_image_left' + '/um_road_' + num + '.png'
187             mask = io.imread(mask_full_path, as_gray=True)
188             mask = cv2.resize(mask, dsize=(im_width, im_height), interpolation=cv2.INTER_CUBIC)
189             mask = mask.reshape((im_height, im_width, 1))
190             masks.append(mask)
191
192     divisor = int(len(images) * split)
193
194     train_data, train_label = np.array(images[:divisor]), np.array(masks[:divisor])
195     test_data, test_label = np.array(images[divisor:]), np.array(masks[divisor:])
196
197     return train_data, train_label, test_data, test_label

```

```

199 if __name__ == '__main__':
200     input_img = Input((im_height, im_width, 1), name='img')
201     model = get_unet(input_img, n_filters=16, dropout=0.05, batchnorm=True)
202
203     model.compile(optimizer=Adam(), loss="binary_crossentropy", metrics=["accuracy"])
204
205     train_data, train_label, test_data, test_label = load_data()
206     results = model.fit(train_data, train_label, batch_size=15, epochs=40, callbacks=callbacks,
207                        validation_data=(test_data, test_label))
208     model.save("road_model_final.h5")
209
210     # Visualize model performance
211     visualize_performance(results)
212
213     # Visualize predict
214     visualize_predict('road_model_final.h5')
215

```

road_visualization.py

```
1 import numpy as np
2 from matplotlib import pyplot as plt
3 from numpy.linalg import lstsq
4 import skimage.io
5 import open3d as o3d
6 from mpl_toolkits import mplot3d
7
8
9 def compute_3d_location(depth, px, py, f):
10     """
11     Project 2: Question 4 & 5
12     Return a matrix of size (h x w x 3) such that  $m[y, x] = [X, Y, Z]$ 
13
14     :param depth: depth map of the stereo image pair
15     :type depth: numpy array
16     :param px: principal point x of camera
17     :type px: float
18     :param py: principal point y of camera
19     :type py: float
20     :param f: focal length
21     :type f: float
22     :return: matrix of the XYZ coordinates in world coordinate system for each pixel
23     :rtype: numpy array
24     """
25
26     xy_coord = make_coordinates_matrix(depth.shape)
27     z_coord = depth.reshape((depth.shape[0], depth.shape[1], 1))
28
29     # calculate X Y coordinates:  $Z * (x - px) / f$ 
30     #  $Z * (y - py) / f$ 
31     xy_coord = (xy_coord - [px, py]) / f
32     xy_coord = np.multiply(xy_coord, np.concatenate((z_coord, z_coord), axis=2))
33
34     location = np.concatenate((xy_coord, z_coord), axis=2)
35     return location
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74 def compute_distance(points, plane):
75     """
76     Helper for RANSAC
77     return a 1d array of size N (number of points) containing the distance of a point to the given plane
78     """
79     # convert 3D point to Homogenous coordinate
80     points = np.concatenate((points, np.ones((points.shape[0], 1))), axis=1)
81     # calculate distance =  $|Ax + By + Cz + d| / \sqrt{A^2 + B^2 + C^2}$ 
82     dists = np.abs(np.dot(points, plane)) / np.sqrt(np.sum(plane[:3] ** 2))
83     return dists
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
```

```

85 def RANSAC(locations, iters=3000, thresh=0.02):
86     """
87     Project 2: Qestion 4
88     RANSAC to estimate a plane robust to outliers
89     :param locations: N x 3 3D points
90     :type locations: numpy array
91     :param iters: number of iterations for RANSAC
92     :type iters: int
93     :param inlier_thresh: point-to-plane distance threshold
94     :type inlier_thresh: float
95     :return: 4D plane vector and a list of inlier indices in the points
96     :rtype:
97     """
98     max_inlier_num = -1
99     best_inlier_idxs = None
100
101     # number of points
102     N = locations.shape[0]
103
104     for i in range(iters):
105         # randomly generate 5 points to fit plane
106         random_idx = np.random.choice(N, 5, replace=False)
107         random_points = locations[random_idx, :]
108
109         # fit plane
110         curr_plane = fit_plane(random_points)
111
112         # calculate distance and find inliers
113         dists = compute_distance(locations, curr_plane)
114         curr_inlier_idxs = np.where(dists < thresh)[0]
115         inlier_num = curr_inlier_idxs.shape[0]
116
117         # update best match
118         if inlier_num > max_inlier_num:
119             max_inlier_num = inlier_num
120             best_inlier_idxs = curr_inlier_idxs
121
122     # using the best match we find so far to fit the final plane
123     best_points = locations[best_inlier_idxs, :]
124     final_plane = fit_plane(best_points)
125
126     # final filter for inliers
127     dists = compute_distance(locations, final_plane)
128     inlier_list = np.where(dists < thresh)[0]
129
130     # [a,b,c,d] -> [a,b,-1,d]
131     final_plane = final_plane / -final_plane[2]
132     return final_plane, inlier_list

```

```

38 def make_coordinates_matrix(im_shape):
39     """
40     Helper for compute_3D_location
41     Return a matrix of size (im_shape[0] x im_shape[1] x 2) such that m[y, x] = [y, x]
42     :param im_shape: (y, x, channel)
43     :type im_shape: tuple
44     :return: xy coordinate matrix of the image
45     :rtype: numpy array
46     """
47     range_x = np.arange(0, im_shape[1], 1)
48     range_y = np.arange(0, im_shape[0], 1)
49     axis_x = np.repeat(range_x[np.newaxis, ...], len(range_y), axis=0)
50     axis_y = np.repeat(range_y[...], np.newaxis, len(range_x), axis=1)
51
52     return np.dstack((axis_y, axis_x))
53
54 def fit_plane(road_loc):
55     """
56     Project 2: Question 4
57     Fit a plane based on the 3D road point cloud
58     :param road_loc: a matrix of size (num_road_pixel x 3)
59     :type road_loc: numpy array
60     :return:
61     :rtype:
62     """
63     # use LSE (least square estimation) to find the best fitting
64     # num_road_pixel x 3 where A[i] = [Xi, Yi, 1]
65     A = np.c_[road_loc[:, 0], road_loc[:, 1], np.ones(road_loc.shape[0])]
66     # run lse to get coefficients
67     C, _, _, _ = lstsq(A, road_loc[:, 2])
68     plane = (C[0], C[1], -1, C[2])
69     nn = np.linalg.norm(plane)
70     plane = plane / nn
71
72     return plane

```

```

135 def filter_pixel(location, road_mask):
136     """
137     Helper function
138     Filter pixel locations and return only the road pixels
139     """
140     # pre-processing mask for pixel classification
141     road_mask[np.where(road_mask < 0.5)] = 0
142     road_mask[np.where(road_mask >= 0.5)] = 1
143     road_mask = road_mask.astype(int)
144     # filter location points and save only road pixels
145     # num_road_pixel x 3
146     road_loc = np.copy(location[np.where(road_mask == 1)])
147
148     return road_loc
149
150 def visualize_plane(X, Y, Z, road_loc):
151     """
152     Visualize the estimated plane with a road point cloud
153     """
154     road_fig = plt.figure()
155     road_ax = plt.axes(projection='3d')
156     road_ax.plot_surface(Z, Y, X, rstride=1, cstride=1, alpha=0.2)
157     road_ax.scatter(road_loc[:,2], road_loc[:,1], road_loc[:,0], cmap='viridis', marker='.', c=road_loc[:,2])
158
159     road_ax.set_xlabel('z')
160     road_ax.set_ylabel('y')
161     road_ax.set_zlabel('x')
162     plt.title('Road 3D Points with Plane', fontsize=16)
163     plt.show()
164
165 def visualize_all_with_plane(X, Y, Z, location, image):
166     """
167     Project 2: Question 5
168     Visualize all the pixel 3D locations and the estimated plane
169     """
170     XX = X.flatten().reshape((-1,1))
171     YY = Y.flatten().reshape((-1,1))
172     ZZ = Z.flatten().reshape((-1,1))
173     road_plane_coord = np.concatenate((XX, YY, ZZ), axis=1)
174
175     # create point cloud for road plane
176     road_pcd = o3d.geometry.PointCloud()
177     road_pcd.points = o3d.utility.Vector3dVector(road_plane_coord)
178     road_pcd.colors = o3d.utility.Vector3dVector([[1, 0, 0] for i in range(XX.shape[0])])
179
180     # create point cloud for all image pixels
181     pcd = o3d.geometry.PointCloud()
182     pcd.points = o3d.utility.Vector3dVector(location.reshape((-1,3)))
183     pcd.colors = o3d.utility.Vector3dVector(image.reshape((-1,3))/255)
184
185     o3d.visualization.draw_geometries([road_pcd, pcd])

```



```

187 def compute_XYZ(points, plane):
188     """
189     Helper for visualization
190     """
191     # find boundary
192     maxx = np.max(points[:, 0])
193     minx = np.min(points[:, 0])
194     maxy = np.max(points[:, 1])
195     miny = np.min(points[:, 1])
196     # create XY meshgrid
197     X, Y = np.meshgrid(np.arange(minx - 50, maxx + 100, 5), np.arange(miny - 50, maxy + 100, 5))
198     # evaluate Z on XY meshgrid
199     Z = plane[0] * X + plane[1] * Y + plane[3]
200     return X, Y, Z
201
202 if __name__ == "__main__":
203     F = 7.215377000000e+02
204     PX, PY = 6.095593000000e+02, 1.728540000000e+02
205
206     depth = skimage.io.imread('./Output Images/depth_umm_000087.png', as_gray=True)
207     # Q4: Fit plane
208     loc = compute_3d_location(depth, PX, PY, F)
209     mask = skimage.io.imread('./Output Images/road_segmentation_mask_umm_000087.png', as_gray=True) / 255
210     road_loc = filter_pixel(loc, mask)
211     plane, inl = RANSAC(road_loc)
212
213     # Q5: Visualize plane
214     left_img = skimage.io.imread('./data/test/image_left/umm_000087.jpg')
215     inl = road_loc[inl]
216     X, Y, Z = compute_XYZ(inl, plane)
217
218     visualize_all_with_plane(X, Y, Z, loc, left_img)
219     visualize_plane(X, Y, Z, road_loc)

```

car_detection.py

```
1 # ref: https://www.learnopencv.com/faster-r-cnn-object-detection-with-pytorch/
2
3 import cv2
4 import torchvision
5 from PIL import Image
6 from torchvision import transforms as T
7 import matplotlib.pyplot as plt
8
9 model = torchvision.models.detection.fasterrcnn_resnet50_fpn(pretrained=True)
10 model.eval()
11
12 COCO_INSTANCE_CATEGORY_NAMES = [
13     '__background__', 'person', 'bicycle', 'car', 'motorcycle', 'airplane', 'bus',
14     'train', 'truck', 'boat', 'traffic light', 'fire hydrant', 'N/A', 'stop sign',
15     'parking meter', 'bench', 'bird', 'cat', 'dog', 'horse', 'sheep', 'cow',
16     'elephant', 'bear', 'zebra', 'giraffe', 'N/A', 'backpack', 'umbrella', 'N/A', 'N/A',
17     'handbag', 'tie', 'suitcase', 'frisbee', 'skis', 'snowboard', 'sports ball',
18     'kite', 'baseball bat', 'baseball glove', 'skateboard', 'surfboard', 'tennis racket',
19     'bottle', 'N/A', 'wine glass', 'cup', 'fork', 'knife', 'spoon', 'bowl',
20     'banana', 'apple', 'sandwich', 'orange', 'broccoli', 'carrot', 'hot dog', 'pizza',
21     'donut', 'cake', 'chair', 'couch', 'potted plant', 'bed', 'N/A', 'dining table',
22     'N/A', 'N/A', 'toilet', 'N/A', 'tv', 'laptop', 'mouse', 'remote', 'keyboard', 'cell phone',
23     'microwave', 'oven', 'toaster', 'sink', 'refrigerator', 'N/A', 'book',
24     'clock', 'vase', 'scissors', 'teddy bear', 'hair drier', 'toothbrush'
25 ]
26
27 VEHICLE_LIST = ['car', 'bus', 'truck']
28
29
30 def get_prediction(img_path, threshold):
31     img = Image.open(img_path)
32
33     # Defining PyTorch transform and apply the transform to the image
34     transform = T.Compose([T.ToTensor()])
35     img = transform(img)
36
37     # Pass the image to the model
38     pred = model([img])
39
40     # Get the labels, bounding boxes and prediction score
41     pred_class = [COCO_INSTANCE_CATEGORY_NAMES[i] for i in list(pred[0]['labels'].numpy())]
42     pred_boxes = [[(i[0], i[1]), (i[2], i[3])] for i in list(pred[0]['boxes'].detach().numpy())]
43     pred_score = list(pred[0]['scores'].detach().numpy())
44
45     # Get list of index with score greater than threshold
46     pred_t = [pred_score.index(x) for x in pred_score if x > threshold][-1]
47     pred_boxes = pred_boxes[:pred_t + 1]
48     pred_class = pred_class[:pred_t + 1]
49
50     return pred_boxes, pred_class
```

```

53 def object_detection_api(img_path, threshold=0.5, rect_th=2, text_size=1, text_th=2):
54     # Get predictions
55     boxes, pred_cls = get_prediction(img_path, threshold)
56
57     img = cv2.imread(img_path)
58     img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
59
60     # Store only vehicle bounding boxes
61     vehicle_box = []
62
63     # Visualize bounding boxes and labels only for vehicles
64     for i in range(len(boxes)):
65         if pred_cls[i] in VIHICLE_LIST:
66             vehicle_box.append(boxes[i])
67             cv2.rectangle(img, boxes[i][0], boxes[i][1], color=(0, 255, 0), thickness=rect_th)
68             cv2.putText(img, pred_cls[i], boxes[i][0], cv2.FONT_HERSHEY_SIMPLEX, text_size, (0, 255, 0),
69                         thickness=text_th)
70
71     # display the output image
72     plt.figure(figsize=(20, 30))
73     plt.imshow(img)
74     plt.xticks([])
75     plt.yticks([])
76     plt.show()
77
78     return vehicle_box
79
80
81 if __name__ == '__main__':
82     vb = object_detection_api('./data/train/image_left/um_000011.jpg')

```

car_viewpoint.py

```
1  # ref: https://www.learnopencv.com/image-classification-using-convolutional-neural-networks-in-keras/
2
3  import os
4  import cv2
5  import numpy as np
6  from skimage import io
7  import scipy.io as sio
8  import matplotlib.pyplot as plt
9  from keras.models import Sequential
10 from keras.layers import Dense, Conv2D, MaxPooling2D, Dropout, Flatten, BatchNormalization, Activation, GaussianNoise
11
12 BATCH_SIZE = 256
13 EPOCHS = 40
14 VALIDATION_SPLIT = 0.8
15 NUM_CLASSES = 12
16
17
18 def read_files(image_path, label_path):
19     # Read images and labels
20     images = []
21     annotations = []
22     for path in os.listdir(image_path):
23         img_full_path = image_path + '/' + path
24         if "DS_Store" not in img_full_path:
25             img = io.imread(img_full_path)
26             images.append(img)
27
28             num = path.split('.')[0]
29             label_full_path = label_path + '/' + num + '.mat'
30             anno = sio.loadmat(label_full_path)['annotation'][0][0]
31             annotations.append(anno)
32
33     return images, annotations
```

```

36 def get_labels(orient):
37     # Use one hot encoder to represent labels
38     divisor = int(orient // 30)
39     label = [1 if i == divisor else 0 for i in range(12)]
40     return np.array(label)
41
42
43 def filter_cars(annotations, img):
44     classes = annotations[0][0]
45     bboxes = annotations[3]
46     orient = annotations[7]
47     truncated = annotations[4][0]
48     occluded = annotations[8][0]
49
50     # Check if each car is good for training, and return bounding boxes and orientations
51     patches, orients = [], []
52     height, width = [], []
53     for i in range(len(classes)):
54         cla = classes[i][0]
55         trunc = truncated[i][0]
56         occ = occluded[i][0]
57         if cla is not 'DontCare' and trunc <= 0.3 and occ <= 2:
58             left, top = int(bboxes[i][0]), int(bboxes[i][1])
59             right, bottom = int(left + bboxes[i][2]), int(top + bboxes[i][3])
60             patch = img[top:bottom, left:right]
61             patches.append(patch)
62             orients.append(get_labels(orient[i][0]))
63             width.append(bboxes[i][2])
64             height.append(bboxes[i][3])
65
66     return height, width, patches, orients

```

```

69 def load_data(image_path, label_path, split):
70     # Read in images and labels
71     images, annotations = read_files(image_path, label_path)
72
73     # Store patch of car and its viewpoint
74     patches_list, labels_list = [], []
75     heights, widths = [], []
76     for i in range(len(images)):
77         height, width, patches, orients = filter_cars(annotations[i], images[i])
78         patches_list.extend(patches)
79         labels_list.extend(orients)
80         heights.extend(height)
81         widths.extend(width)
82
83     # Calculate average patch width and height
84     avg_height = int(sum(heights) / len(heights))
85     avg_width = int(sum(widths) / len(widths))
86
87     # Resize patches to have same size
88     patches_list = [cv2.resize(patch, dsize=(avg_height, avg_width), interpolation=cv2.INTER_CUBIC) for patch in
89                     patches_list]
90
91     # Divide the data into train and test sets with split
92     divider = int(len(patches_list) * split)
93     train_data, train_labels = np.array(patches_list[:divider]), np.array(labels_list[:divider])
94     test_data, test_labels = np.array(patches_list[divider:]), np.array(labels_list[divider:])
95
96     return avg_height, avg_width, train_data, train_labels, test_data, test_labels
97

```

```

99 def createModel(input_shape, nclasses):
100     model = Sequential()
101     model.add(Conv2D(32, (3, 3), padding='same', activation='relu', input_shape=input_shape))
102     model.add(Conv2D(32, (3, 3), activation='relu'))
103     model.add(MaxPooling2D(pool_size=(2, 2)))
104     model.add(Dropout(0.25))
105
106     model.add(Conv2D(64, (3, 3), padding='same', activation='relu'))
107     model.add(Conv2D(64, (3, 3), activation='relu'))
108     model.add(MaxPooling2D(pool_size=(2, 2)))
109     model.add(Dropout(0.25))
110
111     model.add(Conv2D(64, (3, 3), padding='same', activation='relu'))
112     model.add(Conv2D(64, (3, 3), activation='relu'))
113     model.add(MaxPooling2D(pool_size=(2, 2)))
114     model.add(Dropout(0.25))
115
116     model.add(Flatten())
117     model.add(Dense(512, activation='relu'))
118     model.add(Dropout(0.5))
119     model.add(Dense(nclasses, activation='softmax'))
120     return model

```

```

123 def train_model(image_path, label_path):
124     # Load the angle data and split it to train and test sets
125     avg_height, avg_width, train_data, train_labels, test_data, test_labels = load_data(image_path, label_path,
126                                                                                           VALIDATION_SPLIT)
127     # Create a model
128     input_shape = (avg_width, avg_height, 3)
129     model = createModel(input_shape, NUM_CLASSES)
130
131     # Compile and fit the model, evaluate the model with test data
132     model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['accuracy'])
133     history = model.fit(train_data, train_labels, batch_size=BATCH_SIZE, epochs=EPOCHS, verbose=1,
134                        validation_data=(test_data, test_labels))
135     model.evaluate(test_data, test_labels)
136
137     # Save model
138     model.save("model.h5")
139
140     return history

```

```

143 def visualize_performance(history):
144     # Loss Curves
145     plt.figure(figsize=[8, 6])
146     plt.plot(history.history['loss'], 'r', linewidth=3.0)
147     plt.plot(history.history['val_loss'], 'b', linewidth=3.0)
148     plt.legend(['Training loss', 'Validation Loss'], fontsize=18)
149     plt.xlabel('Epochs ', fontsize=16)
150     plt.ylabel('Loss', fontsize=16)
151     plt.title('Loss Curves', fontsize=16)
152     plt.show()
153
154     # Accuracy Curves
155     plt.figure(figsize=[8, 6])
156     plt.plot(history.history['accuracy'], 'r', linewidth=3.0)
157     plt.plot(history.history['val_accuracy'], 'b', linewidth=3.0)
158     plt.legend(['Training Accuracy', 'Validation Accuracy'], fontsize=18)
159     plt.xlabel('Epochs ', fontsize=16)
160     plt.ylabel('Accuracy', fontsize=16)
161     plt.title('Accuracy Curves', fontsize=16)
162     plt.show()
163
164
165 if __name__ == '__main__':
166     history = train_model('data/train_angle/image', 'data/train_angle/labels')
167     visualize_performance(history)

```

car_visualization.py

```
1 import cv2
2 import matplotlib.pyplot as plt
3 import numpy as np
4 from Project.car_detection import object_detection_api
5 from keras.engine.saving import load_model
6 from skimage import io
7
8 ANGLE_DICT = {0: (0, -2), 1: (-1, -np.sqrt(3)), 2: (-np.sqrt(3), -1), 3: (-2, 0), 4: (-np.sqrt(3), 1),
9              5: (-1, np.sqrt(3)), 6: (0, 2), 7: (1, np.sqrt(3)), 8: (np.sqrt(3), 1), 9: (2, 0),
10             10: (np.sqrt(3), -1), 11: (1, -np.sqrt(3))}
11
12 ANGLE_LABEL = {0: '0°', 1: '30°', 2: '60°', 3: '90°', 4: '120°', 5: '150°', 6: '180°', 7: '210°', 8: '240°', 9: '270°',
13              10: '300°', 11: '330°'}
14
15
16 def visualize_car(img_path):
17     # Get car detection result and trained model
18     boxes = object_detection_api(img_path)
19     model = load_model('model.h5')
20
21     img = io.imread(img_path)
22     fig, ax = plt.subplots()
23     ax.imshow(img)
24
25     for i in range(len(boxes)):
26         # Use bounding boxes to crop patch
27         left, top = int(boxes[i][0][0]), int(boxes[i][0][1])
28         right, bottom = int(boxes[i][1][0]), int(boxes[i][1][1])
29
30         # Ignore patches that are too small
31         width, height = right - left, bottom - top
32         if width >= 30 and height >= 30:
33             patch = img[top:bottom, left:right]
34             patch = cv2.resize(patch, dsize=(97, 53), interpolation=cv2.INTER_CUBIC)
35             patch = patch.reshape((1, 97, 53, 3))
36
37             # Predict viewpoint
38             prediction = model.predict_classes(patch)
39             dx, dy = ANGLE_DICT[prediction[0]][0], ANGLE_DICT[prediction[0]][1]
40
41             # Visualize car bounding box and angle
42             center_x, center_y = int((left + right) / 2), int((top + bottom) / 2)
43             rect = plt.Rectangle((left, top), width, height, fill=False, linewidth=2, color='lime')
44             plt.arrow(center_x, center_y, dx * 20, -dy * 20, color='w', linewidth=3, head_width=20, head_length=4)
45             plt.text(center_x, center_y, ANGLE_LABEL[prediction[0]], fontsize=12, color='lime')
46             ax.add_patch(rect)
47
48     plt.show()
49
50
51 if __name__ == '__main__':
52     visualize_car('./data/train/image_left/um_000011.jpg')
```